

Copyright 2025 Google LLC.

```
[ ]: # @title Licensed under the Apache License, Version 2.0 (the "License");
# you may not use this file except in compliance with the License.
# You may obtain a copy of the License at
#
# https://www.apache.org/licenses/LICENSE-2.0
#
# Unless required by applicable law or agreed to in writing, software
# distributed under the License is distributed on an "AS IS" BASIS,
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.
# See the License for the specific language governing permissions and
# limitations under the License.
```

## 🚀 Multi-Agent Systems & Workflow Patterns

Welcome to the Kaggle 5-day Agents course!

In the previous notebook, you built a **single agent** that could take action. Now, you'll learn how to scale up by building **agent teams**.

Just like a team of people, you can create specialized agents that collaborate to solve complex problems. This is called a **multi-agent system**, and it's one of the most powerful concepts in AI agent development.

In this notebook, you'll:

- Learn when to use multi-agent systems in [Agent Development Kit \(ADK\)](#)
- Build your first system using an LLM as a "manager"
- Learn three core workflow patterns (Sequential, Parallel, and Loop) to coordinate your agent teams

### !! Please Read

  Note: No submission required! This notebook is for your hands-on practice and learning only. You do not need to submit it anywhere to complete the course.

 Note: When you first start the notebook via running a cell you might see a banner in the notebook header that reads "Waiting for the next available notebook". The queue should drop rapidly; however, during peak bursts you might have to wait a few minutes.

 Note: Avoid using the Run all cells command as this can trigger a QPM limit resulting in 429 errors when calling the backing model. Suggested flow is to run each cell in order - one at a time. See FAQ on 429 errors for more information.

For help: Ask questions on the [Kaggle Discord server](#).

## 📘 Get started with Kaggle Notebooks

If this is your first time using Kaggle Notebooks, welcome! You can learn more about using Kaggle Notebooks in the documentation.

Here's how to get started:

### 1. Verify Your Account (Required)

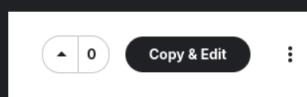
To use the Kaggle Notebooks in this course, you'll need to verify your account with a phone number.

You can do this in your [Kaggle settings](#).

### 2. Make Your Own Copy

To run any code in this notebook, you first need your own editable copy.

Click the [Copy](#) and [Edit](#) button in the top-right corner.

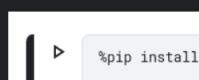


This creates a private copy of the notebook just for you.

### 3. Run Code Cells

Once you have your copy, you can run code.

Click the  Run button next to any code cell to execute it.



Run the cells in order from top to bottom.

### 4. If You Get Stuck

To restart: Select [Factory reset](#) from the [Run](#) menu.

For help: Ask questions on the [Kaggle Discord server](#).

## ⚙️ Section 1: Setup

### 1.1: Install dependencies

This Kaggle Notebooks environment includes a pre-installed version of the `google-adk` library for Python and its required dependencies, so you don't need to install additional packages in this notebook.

To install and use ADK in your own Python development environment outside of this course, you can do so by running:

```
pip install google-adk
```

### 1.2: Configure your Gemini API Key

This notebook uses the [Gemini API](#), which requires authentication.

#### 1. Get your API key

If you don't have one already, create an [API key in Google AI Studio](#).

#### 2. Add the key to Kaggle Secrets

Next, you will need to add your API key to your Kaggle Notebook as a Kaggle User Secret.

1. In the top menu bar of the notebook editor, select `Add-ons` then `Secrets`.
2. Create a new secret with the label `GOOGLE_API_KEY`.
3. Paste your API key into the "Value" field and click "Save".
4. Ensure that the checkbox next to `GOOGLE_API_KEY` is selected so that the secret is attached to the notebook.

#### 3. Authenticate in the notebook

Run the cell below to complete authentication.

```
[1]:  
import os  
from kaggle_secrets import UserSecretsClient  
  
try:  
    GOOGLE_API_KEY = UserSecretsClient().get_secret("GOOGLE_API_KEY")  
    os.environ["GOOGLE_API_KEY"] = GOOGLE_API_KEY  
    print("✅ Gemini API key setup complete.")  
except Exception as e:  
    print(  
        f"⚠️ Authentication Error: Please make sure you have added 'GOOGLE_API_KEY' to your Kaggle secrets. Details: {e}"  
    )  
  
✅ Gemini API key setup complete.
```

### 1.3: Import ADK components

Now, import the specific components you'll need from the Agent Development Kit and the Generative AI library. This keeps your code organized and ensures we have access to the necessary building blocks.

```
[2]:  
from google.adk.agents import Agent, SequentialAgent, ParallelAgent, LoopAgent  
from google.adk.models.google_llm import Gemini  
from google.adk.runners import InMemoryRunner  
from google.adk.tools import AgentTool, FunctionTool, google_search  
from google.genai import types  
  
print("✅ ADK components imported successfully.")  
  
✅ ADK components imported successfully.
```

### 1.4: Configure Retry Options

When working with LLMs, you may encounter transient errors like rate limits or temporary service unavailability. Retry options automatically handle these failures by retrying the request with exponential backoff.

```
[5]:  
retry_config=types.HttpRetryOptions(  
    attempts=5, # Maximum retry attempts  
    exp_base=2, # Delay multiplier  
    initial_delay=1,  
    http_status_codes=[429, 500, 503, 504], # Retry on these HTTP errors  
)
```

## 🧐 Section 2: Why Multi-Agent Systems? + Your First Multi-Agent

### The Problem: The "Do-It-All" Agent

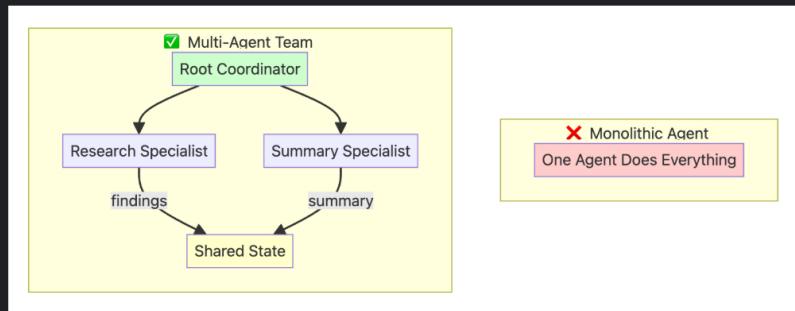
Single agents can do a lot. But what happens when the task gets complex? A single "monolithic" agent that tries to do research, writing, editing, and fact-checking all at once becomes a problem. Its instruction prompt gets long and confusing. It's hard to debug (which part failed?), difficult to maintain, and often produces unreliable results.

### The Solution: A Team of Specialists

Instead of one "do-it-all" agent, we can build a **multi-agent system**. This is a team of simple, specialized agents that collaborate, just like a real-world team. Each agent has one clear job (e.g., one agent *only* does research, another *only* writes). This makes them easier to build, easier to test, and much more powerful and reliable when working together.

To learn more, check out the documentation related to LLM agents in ADK.

#### Architecture: Single Agent vs Multi-Agent Team



## 2.1 Example: Research & Summarization System

Let's build a system with two specialized agents:

1. **Research Agent** - Searches for information using Google Search
2. **Summarizer Agent** - Creates concise summaries from research findings

```
[6]: # Research Agent: Its job is to use the google_search tool and present findings.
research_agent = Agent(
    name="ResearchAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""You are a specialized research agent. Your only job is to use the
google_search tool to find 2-3 pieces of relevant information on the given topic and present the findings with citations.""",
    tools=[google_search],
    output_key="research_findings", # The result of this agent will be stored in the session state with this key.
)
print("✓ research_agent created.")

✓ research_agent created.
```

```
[7]: # Summarizer Agent: Its job is to summarize the text it receives.
summarizer_agent = Agent(
    name="SummarizerAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    # The instruction is modified to request a bulleted list for a clear output format.
    instruction="""Read the provided research findings: {research_findings}
Create a concise summary as a bulleted list with 3-5 key points."""
    output_key="final_summary",
)
print("✓ summarizer_agent created.")

✓ summarizer_agent created.
```

Refer to the ADK documentation for more information on [guiding agents with clear and specific instructions](#).

Then we bring the agents together under a root agent, or coordinator:

```
[8]: # Root Coordinator: Orchestrates the workflow by calling the sub-agents as tools.
root_agent = Agent(
    name="ResearchCoordinator",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    # This instruction tells the root agent HOW to use its tools (which are the other agents).
    instruction="""You are a research coordinator. Your goal is to answer the user's query by orchestrating a workflow.
1. First, you MUST call the 'ResearchAgent' tool to find relevant information on the topic provided by the user.
2. Next, after receiving the research findings, you MUST call the 'SummarizerAgent' tool to create a concise summary.
3. Finally, present the final summary clearly to the user as your response."""
    tools=[AgentTool(research_agent), AgentTool(summarizer_agent)],
)
print("✓ root_agent created.")

✓ root_agent created.
```

Here we're using `AgentTool` to wrap the sub-agents to make them callable tools for the root agent. We'll explore `AgentTool` in detail on Day 2.

Let's run the agent and ask it about a topic:

```
[10]: runner = InMemoryRunner(agent=root_agent)
response = await runner.run_debug(
    "What are the latest advancements in quantum computing and what do they mean for AI?"
)

### Created new session: debug_session_id

User > What are the latest advancements in quantum computing and what do they mean for AI?
WARNING:google_genai.types:Warning: there are non-text parts in the response: ['function_call'], returning concatenated text result from text parts. check the full candidates.content.parts accessor to get the full model response.
WARNING:google_genai.types:Warning: there are non-text parts in the response: ['function_call'], returning concatenated text result from text parts. check the full candidates.content.parts accessor to get the full model response.
ResearchCoordinator > Recent advancements in quantum computing are poised to revolutionize AI by overcoming classical computing limitations and enabling "Quantum AI." This integration promises significantly faster AI training and performance, and the capacity to solve complex, previously intractable problems in fields like drug discovery, logistics, and materials science. Quantum AI also offers more efficient data processing and can lead to more accurate AI models, improving areas such as natural language processing, cybersecurity, and autonomous systems. The integration is expected to drive major innovations across diverse sectors including healthcare, mobility, finance, and energy. While widespread adoption is still some years off, progress is rapid, with increasing accessibility through cloud services, positioning quantum computing as a strategic accelerator for AI.
```

You've just built your first multi-agent system! You used a single "coordinator" agent to manage the workflow, which is a powerful and flexible pattern.

!! However, relying on an LLM's instructions to control the order can sometimes be unpredictable. Next, we'll explore a different pattern that gives you guaranteed, step-by-step execution.

## • Section 3: Sequential Workflows - The Assembly Line

### The Problem: Unpredictable Order

The previous multi-agent system worked, but it relied on a **detailed instruction prompt** to force the LLM to run steps in order. This can be unreliable. A complex LLM might decide to skip a step, run them in the wrong order, or get "stuck," making the process unpredictable.

### The Solution: A Fixed Pipeline

When you need tasks to happen in a **guaranteed, specific order**, you can use a `SequentialAgent`. This agent acts like an assembly line, running each sub-agent in the exact order you list them. The output of one agent automatically becomes the input for the next, creating a predictable and reliable workflow.

**Use Sequential when:** Order matters, you need a linear pipeline, or each step builds on the previous one.

To learn more, check out the documentation related to [sequential agents in ADK](#).

### Architecture: Blog Post Creation Pipeline



## 3.1 Example: Blog Post Creation with Sequential Agents

Let's build a system with three specialized agents:

1. **Outline Agent** - Creates a blog outline for a given topic
2. **Writer Agent** - Writes a blog post
3. **Editor Agent** - Edits a blog post draft for clarity and structure

```
[11]: # Outline Agent: Creates the initial blog post outline.
outline_agent = Agent(
    name="OutlineAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""Create a blog outline for the given topic with:
1. A catchy headline
2. An introduction hook
3. 3-5 main sections with 2-3 bullet points for each
4. A concluding thought""",
    output_key="blog_outline", # The result of this agent will be stored in the session state with this key.
)
print("✓ outline_agent created.")

✓ outline_agent created.
```

```
[12]: # Writer Agent: Writes the full blog post based on the outline from the previous agent.
writer_agent = Agent(
    name="WriterAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""Write a detailed blog post based on the outline provided. The outline is a list of sections with bullet points. Your response should be a well-structured blog post that includes an introduction, several main sections, and a conclusion. Use the same structure as the outline, including headings and bullet points where applicable.
{outline}""",
    output_key="blog_post"
)
print("✓ writer_agent created.")

✓ writer_agent created.
```

```

),
    # The '{blog_outline}' placeholder automatically injects the state value from the previous agent's output.
    instruction="""Following this outline strictly: {blog_outline}
    Write a brief, 200 to 300-word blog post with an engaging and informative tone."""
    output_key="blog_draft", # The result of this agent will be stored with this key.
)

print("✅ writer_agent created.")

✓ writer_agent created.

```

```

[13]: # Editor Agent: Edits and polishes the draft from the writer agent.
editor_agent = Agent(
    name="EditorAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    # This agent receives the '{blog_draft}' from the writer agent's output.
    instruction="""Edit this draft: {blog_draft}
    Your task is to polish the text by fixing any grammatical errors, improving the flow and sentence structure, and enhancing overall clarity."""
    output_key="final_blog", # This is the final output of the entire pipeline.
)

print("✅ editor_agent created.")

✓ editor_agent created.

```

Then we bring the agents together under a sequential agent, which runs the agents in the order that they are listed:

```

[14]: root_agent = SequentialAgent(
    name="BlogPipeline",
    sub_agents=[outline_agent, writer_agent, editor_agent],
)

print("✅ Sequential Agent created.")

✓ Sequential Agent created.

```

Let's run the agent and give it a topic to write a blog post about:

```

▶ runner = InMemoryRunner(agent=root_agent)
response = await runner.run_debug(
    "Write a blog post about the benefits of multi-agent systems for software developers"
)

### Created new session: debug_session_id

User > Write a blog post about the benefits of multi-agent systems for software developers
OutlineAgent > ## Blog outline:

**Headline:** Unleash Your Development Superpowers: How Multi-Agent Systems Are Revolutionizing Software

**Introduction Hook:** Tired of monolithic codebases that buckle under complexity? Imagine a software architecture where independent, intelligent agents collaborate to achieve a common goal, making your development process more agile, scalable, and robust. Welcome to the world of Multi-Agent Systems (MAS), a paradigm shift poised to transform how we build software.

---

### Main Sections:

**1. Enhanced Modularity and Maintainability:**

* **Breaking Down the Monolith:** MAS allows developers to decompose complex problems into smaller, more manageable units (agents). This promotes clear separation of concerns, making individual components easier to understand, debug, and modify.
* **Independent Evolution:** Each agent can be developed, tested, and updated independently. This drastically reduces the risk of introducing unintended side effects and accelerates the development lifecycle.
* **Simplified Troubleshooting:** When issues arise, developers can pinpoint problems to specific agents, rather than sifting through vast amounts of interconnected code.

**2. Improved Scalability and Resilience:**

* **Dynamic Scaling:** MAS naturally supports scalability. As demand increases, new agents can be instantiated and deployed to handle the load without disrupting existing functionality.
* **Fault Tolerance:** If one agent fails, others can often compensate or take over its responsibilities, leading to more resilient systems that can withstand partial failures.
* **Distributed Processing:** Agents can operate on different machines or even different geographical locations, enabling true distributed computing and leveraging available resources more effectively.

**3. Fostering Collaboration and Intelligence:**

* **Complex Problem Solving:** MAS enables the tackling of highly complex problems that might be intractable for a single, monolithic system. Agents can share information, negotiate, and coordinate their actions to find optimal solutions.
* **Emergent Behavior:** The interaction of simple agents can lead to sophisticated, emergent behaviors that were not explicitly programmed, unlocking new levels of intelligence and adaptability in your software.
* **Adaptability to Changing Environments:** Agents can be designed to sense their environment and adapt their behavior accordingly, making MAS ideal for dynamic and unpredictable situations.

---

**Concluding Thought:** Embracing Multi-Agent Systems isn't just about adopting a new architectural pattern; it's about unlocking a more powerful, flexible, and intelligent way to develop software. For developers, it means wielding tools that amplify their ability to build the complex, scalable, and resilient applications of the future. Are you ready to join the agent revolution?

```

```
ClientError                                     Traceback (most recent call last)
/tmp/ipykernel_48/1463792725.py in <cell line: 1>()
      1 runner = InMemoryRunner(agent=root_agent)
----> 2 response = await runner.run_debug(
      3     "Write a blog post about the benefits of multi-agent systems for software developers"
      4 )

/usr/local/lib/python3.11/dist-packages/google/adk/runners.py in run_debug(self, user_messages, user_id, session_id, run_config, quiet, verbose)
  1021     print(f'\nUser > {message}')
  1022
-> 1023     async for event in self.run_async(
  1024         user_id=user_id,
  1025         session_id=session_id,
  1026     )
  1027
  1028     # Run compaction after all events are yielded from the agent.

/usr/local/lib/python3.11/dist-packages/google/adk/runners.py in run_async(self, user_id, session_id, invocation_id, new_message, state_delta, run_config)
  441
  442     async with Aclosing(_run_with_trace(new_message, invocation_id)) as agen:
--> 443         async for event in agen:
  444             yield event
  445
  446
  447     # Run compaction after all events are yielded from the agent.

/usr/local/lib/python3.11/dist-packages/google/adk/runners.py in _run_with_trace(new_message, invocation_id)
  425
  426     ) as agen:
--> 427         async for event in agen:
  428             yield event
  429
  430     # Run compaction after all events are yielded from the agent.

/usr/local/lib/python3.11/dist-packages/google/adk/runners.py in _exec_with_plugin(self, invocation_context, session, execute_fn, is_live_call)
  651     # Step 2: Otherwise continue with normal execution
  652     async with Aclosing(execute_fn(invocation_context)) as agen:
--> 653         async for event in agen:
  654             if not event.partial:
  655                 if self._should_append_event(event, is_live_call):
  656
  657
  658     # Step 3: Continue with normal execution
  659     # ...
  660
  661     # Step 4: Continue with normal execution
  662     # ...
  663
  664     # Step 5: Continue with normal execution
  665     # ...
  666
  667     # Step 6: Continue with normal execution
  668     # ...
  669
  670     # Step 7: Continue with normal execution
  671     # ...
  672
  673     # Step 8: Continue with normal execution
  674     # ...
  675
  676     # Step 9: Continue with normal execution
  677     # ...
  678
  679     # Step 10: Continue with normal execution
  680     # ...
  681
  682     # Step 11: Continue with normal execution
  683     # ...
  684
  685     # Step 12: Continue with normal execution
  686     # ...
  687
  688     # Step 13: Continue with normal execution
  689     # ...
  690
  691     # Step 14: Continue with normal execution
  692     # ...
  693
  694     # Step 15: Continue with normal execution
  695     # ...
  696
  697     # Step 16: Continue with normal execution
  698     # ...
  699
  700     # Step 17: Continue with normal execution
  701     # ...
  702
  703     # Step 18: Continue with normal execution
  704     # ...
  705
  706     # Step 19: Continue with normal execution
  707     # ...
  708
  709     # Step 20: Continue with normal execution
  710     # ...
  711
  712     # Step 21: Continue with normal execution
  713     # ...
  714
  715     # Step 22: Continue with normal execution
  716     # ...
  717
  718     # Step 23: Continue with normal execution
  719     # ...
  720
  721     # Step 24: Continue with normal execution
  722     # ...
  723
  724     # Step 25: Continue with normal execution
  725     # ...
  726
  727     # Step 26: Continue with normal execution
  728     # ...
  729
  730     # Step 27: Continue with normal execution
  731     # ...
  732
  733     # Step 28: Continue with normal execution
  734     # ...
  735
  736     # Step 29: Continue with normal execution
  737     # ...
  738
  739     # Step 30: Continue with normal execution
  740     # ...
  741
  742     # Step 31: Continue with normal execution
  743     # ...
  744
  745     # Step 32: Continue with normal execution
  746     # ...
  747
  748     # Step 33: Continue with normal execution
  749     # ...
  750
  751     # Step 34: Continue with normal execution
  752     # ...
  753
  754     # Step 35: Continue with normal execution
  755     # ...
  756
  757     # Step 36: Continue with normal execution
  758     # ...
  759
  760     # Step 37: Continue with normal execution
  761     # ...
  762
  763     # Step 38: Continue with normal execution
  764     # ...
  765
  766     # Step 39: Continue with normal execution
  767     # ...
  768
  769     # Step 40: Continue with normal execution
  770     # ...
  771
  772     # Step 41: Continue with normal execution
  773     # ...
  774
  775     # Step 42: Continue with normal execution
  776     # ...
  777
  778     # Step 43: Continue with normal execution
  779     # ...
  780
  781     # Step 44: Continue with normal execution
  782     # ...
  783
  784     # Step 45: Continue with normal execution
  785     # ...
  786
  787     # Step 46: Continue with normal execution
  788     # ...
  789
  790     # Step 47: Continue with normal execution
  791     # ...
  792
  793     # Step 48: Continue with normal execution
  794     # ...
  795
  796     # Step 49: Continue with normal execution
  797     # ...
  798
  799     # Step 50: Continue with normal execution
  800     # ...
  801
  802     # Step 51: Continue with normal execution
  803     # ...
  804
  805     # Step 52: Continue with normal execution
  806     # ...
  807
  808     # Step 53: Continue with normal execution
  809     # ...
  810
  811     # Step 54: Continue with normal execution
  812     # ...
  813
  814     # Step 55: Continue with normal execution
  815     # ...
  816
  817     # Step 56: Continue with normal execution
  818     # ...
  819
  820     # Step 57: Continue with normal execution
  821     # ...
  822
  823     # Step 58: Continue with normal execution
  824     # ...
  825
  826     # Step 59: Continue with normal execution
  827     # ...
  828
  829     # Step 60: Continue with normal execution
  830     # ...
  831
  832     # Step 61: Continue with normal execution
  833     # ...
  834
  835     # Step 62: Continue with normal execution
  836     # ...
  837
  838     # Step 63: Continue with normal execution
  839     # ...
  840
  841     # Step 64: Continue with normal execution
  842     # ...
  843
  844     # Step 65: Continue with normal execution
  845     # ...
  846
  847     # Step 66: Continue with normal execution
  848     # ...
  849
  850     # Step 67: Continue with normal execution
  851     # ...
  852
  853     # Step 68: Continue with normal execution
  854     # ...
  855
  856     # Step 69: Continue with normal execution
  857     # ...
  858
  859     # Step 70: Continue with normal execution
  860     # ...
  861
  862     # Step 71: Continue with normal execution
  863     # ...
  864
  865     # Step 72: Continue with normal execution
  866     # ...
  867
  868     # Step 73: Continue with normal execution
  869     # ...
  870
  871     # Step 74: Continue with normal execution
  872     # ...
  873
  874     # Step 75: Continue with normal execution
  875     # ...
  876
  877     # Step 76: Continue with normal execution
  878     # ...
  879
  880     # Step 77: Continue with normal execution
  881     # ...
  882
  883     # Step 78: Continue with normal execution
  884     # ...
  885
  886     # Step 79: Continue with normal execution
  887     # ...
  888
  889     # Step 80: Continue with normal execution
  890     # ...
  891
  892     # Step 81: Continue with normal execution
  893     # ...
  894
  895     # Step 82: Continue with normal execution
  896     # ...
  897
  898     # Step 83: Continue with normal execution
  899     # ...
  900
  901     def __get_llm(self, invocation_context: InvocationContext) -> BaseLlm:
  902
  903         try:
  904             with Aclosing(response_generator) as agen:
```

```

    187     yield from self._process_response(response_generator)
--> 188     async for response in agen:
189         yield response
190
191     except Exception as model_error:
192
193         raise model_error
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215
216
217
218
219
220
221
222
223
224
225
226
227
228
229
230
231
232
233
234
235
236
237
238
239
240
241
242
243
244
245
246
247
248
249
250
251
252
253
254
255
256
257
258
259
260
261
262
263
264
265
266
267
268
269
270
271
272
273
274
275
276
277
278
279
280
281
282
283
284
285
286
287
288
289
290
291
292
293
294
295
296
297
298
299
300
301
302
303
304
305
306
307
308
309
310
311
312
313
314
315
316
317
318
319
320
321
322
323
324
325
326
327
328
329
330
331
332
333
334
335
336
337
338
339
340
341
342
343
344
345
346
347
348
349
350
351
352
353
354
355
356
357
358
359
360
361
362
363
364
365
366
367
368
369
370
371
372
373
374
375
376
377
378
379
380
381
382
383
384
385
386
387
388
389
390
391
392
393
394
395
396
397
398
399
400
401
402
403
404
405
406
407
408
409
410
411
412
413
414
415
416
417
418
419
420
421
422
423
424
425
426
427
428
429
430
431
432
433
434
435
436
437
438
439
440
441
442
443
444
445
446
447
448
449
450
451
452
453
454
455
456
457
458
459
460
461
462
463
464
465
466
467
468
469
470
471
472
473
474
475
476
477
478
479
480
481
482
483
484
485
486
487
488
489
490
491
492
493
494
495
496
497
498
499
500
501
502
503
504
505
506
507
508
509
510
511
512
513
514
515
516
517
518
519
520
521
522
523
524
525
526
527
528
529
530
531
532
533
534
535
536
537
538
539
540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593
594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647
648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701
702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755
756
757
758
759
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
779
780
781
782
783
784
785
786
787
788
789
789
790
791
792
793
794
795
796
797
798
799
799
800
801
802
803
804
805
806
807
808
809
809
810
811
812
813
814
815
816
817
818
819
819
820
821
822
823
824
825
826
827
828
829
829
830
831
832
833
834
835
836
837
838
839
839
840
841
842
843
844
845
846
847
847
848
849
849
850
851
852
853
854
855
856
857
858
859
859
860
861
862
863
864
865
866
867
868
869
869
870
871
872
873
874
875
876
877
878
879
879
880
881
882
883
884
885
886
887
888
889
889
890
891
892
893
894
895
896
897
898
899
899
900
901
902
903
904
905
906
907
908
909
909
910
911
912
913
914
915
916
917
918
919
919
920
921
922
923
924
925
926
927
928
929
929
930
931
932
933
934
935
936
937
938
939
939
940
941
942
943
944
945
946
947
947
948
949
949
950
951
952
953
954
955
956
957
958
959
959
960
961
962
963
964
965
966
967
968
969
969
970
971
972
973
974
975
976
977
978
979
979
980
981
982
983
984
985
986
987
988
989
989
990
991
992
993
994
995
996
997
998
999
999
1000
1000
1001
1002
1003
1004
1005
1006
1007
1008
1009
1009
1010
1011
1012
1013
1014
1015
1016
1017
1018
1019
1019
1020
1021
1022
1023
1024
1025
1026
1027
1028
1029
1029
1030
1031
1032
1033
1034
1035
1036
1037
1038
1039
1039
1040
1041
1042
1043
1044
1045
1046
1047
1047
1048
1049
1049
1050
1051
1052
1053
1054
1055
1056
1057
1058
1059
1059
1060
1061
1062
1063
1064
1065
1066
1067
1068
1069
1069
1070
1071
1072
1073
1074
1075
1076
1077
1078
1079
1079
1080
1081
1082
1083
1084
1085
1086
1087
1088
1089
1089
1090
1091
1092
1093
1094
1095
1096
1097
1098
1099
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1139
1140
1141
1142
1143
1144
1145
1146
1147
1147
1148
1149
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1789
1790
1791
1792
1793
1794
1795
1796
1797
1797
1798
1799
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1998
1999
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2098
2099
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2198
2199
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2298
2299
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2398
2399
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2439
2440
244
```

## Section 4: Parallel Workflows - Independent Researchers

### The Problem: The Bottleneck

The previous sequential agent is great, but it's an assembly line. Each step must wait for the previous one to finish. What if you have several tasks that are **not dependent** on each other? For example, researching three *different* topics. Running them in sequence would be slow and inefficient, creating a bottleneck where each task waits unnecessarily.

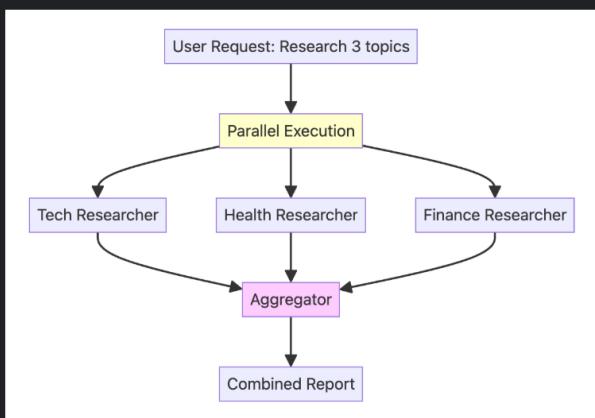
### The Solution: Concurrent Execution

When you have independent tasks, you can run them all at the same time using a `ParallelAgent`. This agent executes all of its sub-agents concurrently, dramatically speeding up the workflow. Once all parallel tasks are complete, you can then pass their combined results to a final 'aggregator' step.

**Use Parallel when:** Tasks are independent, speed matters, and you can execute concurrently.

To learn more, check out the documentation related to parallel agents in ADK.

### Architecture: Multi-Topic Research



### 4.1 Example: Parallel Multi-Topic Research

Let's build a system with four agents:

1. **Tech Researcher** - Researches AI/ML news and trends
2. **Health Researcher** - Researches recent medical news and trends
3. **Finance Researcher** - Researches finance and fintech news and trends
4. **Aggregator Agent** - Combines all research findings into a single summary

```
[16]: # Tech Researcher: Focuses on AI and ML trends.
tech_researcher = Agent(
    name="TechResearcher",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""Research the latest AI/ML trends. Include 3 key developments,
the main companies involved, and the potential impact. Keep the report very concise (100 words).""",
    tools=[google_search],
    output_key="tech_research", # The result of this agent will be stored in the session state with this key.
)

print("✅ tech_researcher created.")

✅ tech_researcher created.
```

```
[17]: # Health Researcher: Focuses on medical breakthroughs.
health_researcher = Agent(
    name="HealthResearcher",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""Research recent medical breakthroughs. Include 3 significant advances,
their practical applications, and estimated timelines. Keep the report concise (100 words).""",
    tools=[google_search],
    output_key="health_research", # The result will be stored with this key.
)

print("✅ health_researcher created.")

✅ health_researcher created.
```

```
[18]: # Finance Researcher: Focuses on fintech trends.
finance_researcher = Agent(
```

```

        name="FinanceResearcher",
        model=Gemini(
            model="gemini-2.5-flash-lite",
            retry_options=retry_config
        ),
        instruction="""Research current fintech trends. Include 3 key trends,
their market implications, and the future outlook. Keep the report concise (100 words).""",
        tools=[google_search],
        output_key="finance_research", # The result will be stored with this key.
    )

    print("✅ finance_researcher created.")

    ✅ finance_researcher created.

```

```

[19]: # The AggregatorAgent runs *after* the parallel step to synthesize the results.
aggregator_agent = Agent(
    name="AggregatorAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    # It uses placeholders to inject the outputs from the parallel agents, which are now in the session state.
    instruction="""Combine these three research findings into a single executive summary:

    **Technology Trends:** {tech_research}

    **Health Breakthroughs:** {health_research}

    **Finance Innovations:** {finance_research}

    Your summary should highlight common themes, surprising connections, and the most important key takeaways from all three reports. The final summary should be a single, cohesive document.""".format(tech_research=tech_research, health_research=health_research, finance_research=finance_research),
    output_key="executive_summary", # This will be the final output of the entire system.
)

print("✅ aggregator_agent created.")

    ✅ aggregator_agent created.

```

👉 Then we bring the agents together under a parallel agent, which is itself nested inside of a sequential agent.

This design ensures that the research agents run first in parallel, then once all of their research is complete, the aggregator agent brings together all of the research findings into a single report:

```

[20]: # The ParallelAgent runs all its sub-agents simultaneously.
parallel_research_team = ParallelAgent(
    name="ParallelResearchTeam",
    sub_agents=[tech_researcher, health_researcher, finance_researcher],
)

# This SequentialAgent defines the high-level workflow: run the parallel team first, then run the aggregator.
root_agent = SequentialAgent(
    name="ResearchSystem",
    sub_agents=[parallel_research_team, aggregator_agent],
)

print("✅ Parallel and Sequential Agents created.")

    ✅ Parallel and Sequential Agents created.

```

Let's run the agent and give it a prompt to research the given topics:

```

[21]: runner = InMemoryRunner(agent=root_agent)
response = await runner.run_debug(
    "Run the daily executive briefing on Tech, Health, and Finance"
)

### Created new session: debug_session_id

User > Run the daily executive briefing on Tech, Health, and Finance
TechResearcher > **Key AI/ML Trends and Developments:**

1. **Generative AI for Complex Content Creation:** AI models are increasingly capable of generating diverse content, including graphics, video, and music, beyond just text. Companies like Google (Imagen, Muse) and OpenAI (GPT models) are at the forefront. This trend enhances creativity and practical applications across industries.

2. **Shift from LLMs to SLMs and Specialized Models:** While Large Language Models (LLMs) have dominated, there's a growing interest in smaller Language Models (SLMs) and highly domain-specific models. This offers efficiency and tailored capabilities.

3. **AI Agents and Autonomous Systems:** AI agents capable of independent action are emerging for tasks like workflow management and data analysis. Companies like Salesforce are developing such tools. This trend promises increased efficiency and automation.

**Main Companies Involved:**

* **Tech Giants:** Google (Alphabet), OpenAI, Microsoft, Meta, Amazon, NVIDIA.
* **Specialized AI Companies:** Anthropic, Synthesia, Runway ML, Dataiku, Mistral AI.

**Potential Impact:**

These AI/ML trends are poised to revolutionize industries by increasing productivity, enhancing decision-making, and fostering innovation. They enable personalized custom...

```

These AI/ML trends are poised to revolutionize industries by increasing productivity, enhancing decision-making, and fostering innovation. They enable personalized customer experiences, automate complex tasks, and drive new avenues for content creation. However, they also bring challenges related to ethical considerations, data privacy, and workforce adaptation, necessitating a balance between technological advancement and responsible implementation.

HealthResearcher > Here's a brief executive briefing on recent breakthroughs in Technology, Health, and Finance:

**Technology:** AI and machine learning continue to advance rapidly, with applications in automation, natural language processing, and personalized experiences. Extended Reality (XR), encompassing VR and AR, is transforming gaming, education, and remote collaboration. Quantum computing is showing significant progress, with potential applications in drug discovery, materials science, and cybersecurity, though widespread practical use is still debated.

**Health:** Next-generation mRNA vaccines offer rapid development and broad applications beyond infectious diseases. Innovative surgical techniques, like the first living mitral valve replacement and total eye/face transplants, are improving patient outcomes. Advancements in AI are also enhancing medical diagnostics and drug discovery.

**Finance:** Embedded finance is integrating financial services into non-financial platforms, projected for substantial growth. Decentralized finance (DeFi) is maturing with increased security and scalability. AI-powered personalization is becoming central, with AI in financial services market rapidly expanding.

FinanceResearcher > Here's a daily executive briefing for November 11, 2025:

**Technology:**

- AI Integration in Enterprises:** Companies like Spark are focusing on making AI accessible to all employees, emphasizing organizational change management as a key challenge. The trend involves balancing short-term gains with long-term AI investments and adapting to evolving AI models.

- Digital Transformation & Data Sovereignty:** Tech conferences are highlighting cloud technologies that respect data sovereignty, particularly in sectors like health care, promoting hybrid-cloud models for secure data management.

- Cybersecurity & Spyware Concerns:** Tech giants are stepping up defenses against spyware, with ongoing litigation and scrutiny of companies involved in surveillance technology.

**Health:**

- Global Health Financing Emergency:** A significant drop in external health aid threatens health systems, especially in low-income countries. There's a push for increased domestic investment and self-reliance in health financing.

- Measles Transmission Resurfaces:** The Americas region has lost its measles-free verification due to reestablished endemic transmission in Canada, highlighting the critical need for high vaccination coverage.

- AI in Digital Mental Health:** The use of AI in digital mental health devices presents both opportunities for improved patient care and risks concerning effectiveness and accessibility, prompting careful regulatory consideration.

**Finance:**

- Climate Finance Mobilization:** Efforts are underway to mobilize private climate finance for adaptation, with a focus on developing new financial models and "resilience asset classes" to bridge the gap between philanthropy and markets.

- Financial Regulation & Open Banking:** Regulatory bodies are actively involved in rulemaking for open banking and stablecoin frameworks, signaling a dynamic regulatory landscape.

- Corporate Earnings & Market Trends:** Companies are reporting on financial performance, with trends including revenue growth from acquisitions, focus on cost savings, and strategic capital management.

AggregatorAgent > ## Executive Summary: Tech, Health, and Finance Briefing

**Key Takeaway:** Artificial Intelligence is the dominant transformative force across technology, health, and finance, driving innovation while presenting significant ethical and logistical challenges.

**Technology:** is rapidly advancing with sophisticated AI, from generative content creation (graphics, video, music) to emerging AI agents for autonomous tasks. The field is shifting towards more efficient Smaller Language Models (SLMs) and specialized AI, alongside continued progress in Extended Reality (XR) and the potential of quantum computing. However, cybersecurity concerns, particularly around spyware, and the need for organizational change management to integrate AI effectively are critical considerations.

In **Health**, AI is revolutionizing diagnostics and drug discovery, complementing breakthroughs in next-generation mRNA vaccines and innovative surgical techniques. Simultaneously, a global health financing emergency, particularly impacting low-income countries, and resurgent infectious diseases like measles underscore the need for robust public health infrastructure and increased domestic investment. The integration of AI in digital mental health also presents a duality of opportunity and risk.

**Finance:** is being reshaped by embedded finance, maturing DeFi, and pervasive AI-powered personalization. A significant focus is on mobilizing climate finance through new models and addressing regulatory landscapes for open banking and stablecoins. Corporate earnings reflect strategic capital management and cost-saving initiatives.

**Connecting Themes:** The consistent thread is the pervasive influence of AI, enhancing capabilities across all sectors while necessitating careful consideration of ethical implications, data sovereignty, accessibility, and the balance between innovation and responsible implementation.

+ Code

+ Markdown

Great! You've seen how parallel agents can dramatically speed up workflows by running independent tasks concurrently.

So far, all our workflows run from start to finish and then stop. **But what if you need to review and improve an output multiple times?** Next, we'll build a workflow that can loop and refine its own work.

## Section 5: Loop Workflows - The Refinement Cycle

### The Problem: One-Shot Quality

All the workflows we've seen so far run from start to finish. The `SequentialAgent` and `ParallelAgent` produce their final output and then stop. This 'one-shot' approach isn't good for tasks that require refinement and quality control. What if the first draft of our story is bad? We have no way to review it and ask for a rewrite.

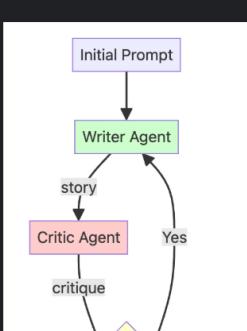
### The Solution: Iterative Refinement

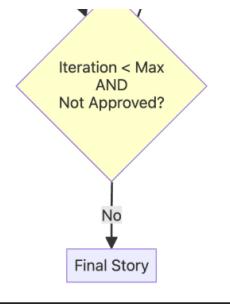
When a task needs to be improved through cycles of feedback and revision, you can use a `LoopAgent`. A `LoopAgent` runs a set of sub-agents repeatedly *until a specific condition is met or a maximum number of iterations is reached*. This creates a refinement cycle, allowing the agent system to improve its own work over and over.

**Use Loop when:** Iterative improvement is needed, quality refinement matters, or you need repeated cycles.

To learn more, check out the documentation related to loop agents in ADK.

### Architecture: Story Writing & Critique Loop





## 5.1 Example: Iterative Story Refinement

Let's build a system with two agents:

1. **Writer Agent** - Writes a draft of a short story
2. **Critic Agent** - Reviews and critiques the short story to suggest improvements

```
[22]: # This agent runs ONCE at the beginning to create the first draft.
initial_writer_agent = Agent(
    name="InitialWriterAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""Based on the user's prompt, write the first draft of a short story (around 100-150 words).
Output only the story text, with no introduction or explanation.""",
    output_key="current_story", # Stores the first draft in the state.
)
print("✓ initial_writer_agent created.")

✓ initial_writer_agent created.
```

```
[23]: # This agent's only job is to provide feedback or the approval signal. It has no tools.
critic_agent = Agent(
    name="CriticAgent",
    model=Gemini(
        model="gemini-2.5-flash-lite",
        retry_options=retry_config
    ),
    instruction="""You are a constructive story critic. Review the story provided below.
Story: {current_story}

Evaluate the story's plot, characters, and pacing.
- If the story is well-written and complete, you MUST respond with the exact phrase: "APPROVED"
- Otherwise, provide 2-3 specific, actionable suggestions for improvement."""
    output_key="critique", # Stores the feedback in the state.
)
print("✓ critic_agent created.")

✓ critic_agent created.
```

Now, we need a way for the loop to actually stop based on the critic's feedback. The `LoopAgent` itself doesn't automatically know that "APPROVED" means "stop."

We need an agent to give it an explicit signal to terminate the loop.

We do this in two parts:

1. A simple Python function that the `LoopAgent` understands as an "exit" signal.
2. An agent that can call that function when the right condition is met.

First, you'll define the `exit_loop` function:

```
[24]: # This is the function that the RefinerAgent will call to exit the loop.
def exit_loop():
    """Call this function ONLY when the critique is 'APPROVED', indicating the story is finished and no more changes are needed."""
    return {"status": "approved", "message": "Story approved. Exiting refinement loop."}

print("✓ exit_loop function created.")

✓ exit_loop function created.
```

To let an agent call this Python function, we wrap it in a `FunctionTool`. Then, we create a `RefinerAgent` that has this tool.

**👉 Notice its instructions:** this agent is the "brain" of the loop. It reads the `{critique}` from the `CriticAgent` and decides whether to (1) call the `exit_loop` tool or (2) rewrite the story.

```
[25]: # This agent refines the story based on critique OR calls the exit_loop function.
refiner_agent = Agent(
    name="RefinerAgent"
```

```

name="refiner_agent",
model=Gemini(
    model="gemini-2.5-flash-lite",
    retry_options=retry_config
),
instruction="""You are a story refiner. You have a story draft and critique.

Story Draft: {current_story}
Critique: {critique}

Your task is to analyze the critique.
- IF the critique is EXACTLY "APPROVED", you MUST call the `exit_loop` function and nothing else.
- OTHERWISE, rewrite the story draft to fully incorporate the feedback from the critique."""",
output_key="current_story", # It overwrites the story with the new, refined version.
tools=[
    FunctionTool(exit_loop)
], # The tool is now correctly initialized with the function reference.

)

print("✅ refiner_agent created.")

✅ refiner_agent created.

```

Then we bring the agents together under a loop agent, which is itself nested inside of a sequential agent.

This design ensures that the system first produces an initial story draft, then the refinement loop runs up to the specified number of `max_iterations`:

```

[26]: # The LoopAgent contains the agents that will run repeatedly: Critic -> Refiner.
story_refinement_loop = LoopAgent(
    name="StoryRefinementLoop",
    sub_agents=[critic_agent, refiner_agent],
    max_iterations=2, # Prevents infinite loops
)

# The root agent is a SequentialAgent that defines the overall workflow: Initial Write -> Refinement Loop.
root_agent = SequentialAgent(
    name="StoryPipeline",
    sub_agents=[initial_writer_agent, story_refinement_loop],
)

print("✅ Loop and Sequential Agents created.")

✅ Loop and Sequential Agents created.

```

Let's run the agent and give it a topic to write a short story about:

```

[28]: runner = InMemoryRunner(agent=root_agent)
response = await runner.run_debug(
    "Write a short story about a lighthouse keeper who discovers a mysterious, glowing map"
)

### Created new session: debug_session_id

User > Write a short story about a lighthouse keeper who discovers a mysterious, glowing map
InitialWriterAgent > Elias traced the salt-worn lines of his map, a familiar comfort against the roaring storm outside. But tonight, something was different. A faint luminescence pulsed from beneath the ancient vellum. He held it closer. The ink, usually faded to near invisibility, now glowed with an ethereal, sea-green light, tracing a path he'd never seen before. It led not to the mainland, but out, into the inky blackness of the churning ocean. A single, glittering X marked a spot far beyond the known shipping lanes. Elias's heart hammered. He'd tended this light for thirty years, seen all manner of things the sea could conjure, but this... this was a secret the ocean had kept for itself, and it was now offering it to him.
CriticAgent > The plot has a strong hook and a clear inciting incident: the discovery of the glowing map. The pacing is effective in building immediate intrigue and mystery.

Here are a few suggestions for improvement:
1. **Develop Elias's character beyond his profession:** While his thirty years as a lighthouse keeper provide context, giving Elias a brief internal thought or a small action that reveals a personal desire or a past longing would make his reaction to this extraordinary discovery more impactful. What has he always *hoped* for, or what has he *lost* that this map might represent?
2. **Hint at the nature of the secret:** The phrase "a secret the ocean had kept for itself" is intriguing, but a slightly more concrete, even if still vague, suggestion could deepen the mystery. Is it treasure? A lost civilization? A mythical creature's lair? A hint, however subtle, would give the reader something more to latch onto.
3. **Show, don't just tell, Elias's emotional response:** You mention his "Heart hammered." Consider adding a physical reaction or a sensory detail to further illustrate his astonishment and apprehension. Does his breath catch? Does he drop the map momentarily? Does the light reflect strangely in his wide eyes?
RefinerAgent > Elias traced the salt-worn lines of his map, a familiar comfort against the roaring storm outside. Thirty years he'd guarded this lonely light, his life a quiet rhythm of sea and sky. But tonight, something was different. A faint luminescence pulsed from beneath the ancient vellum. He held it closer, his weathered hand trembling. The ink, usually faded to near invisibility, now glowed with an ethereal, sea-green light, tracing a path he'd never seen before. It led not to the mainland, but out, into the inky blackness of the churning ocean. A single, glittering X marked a spot far beyond the known shipping lanes. Elias's breath hitched, the map nearly slipping from his grasp as the light reflected in his wide, astonished eyes. He'd always dreamed of more than just the endless horizon, of a purpose beyond the beam he swept each night. Could this be it? This glowing mark, hinting at forgotten depths, perhaps sunken cities or the fabled kraken's hoard – a secret the ocean had guarded for millennia, and was now offering to him. His heart hammered not with fear, but with a wild, unexpected hope.
CriticAgent > APPROVED

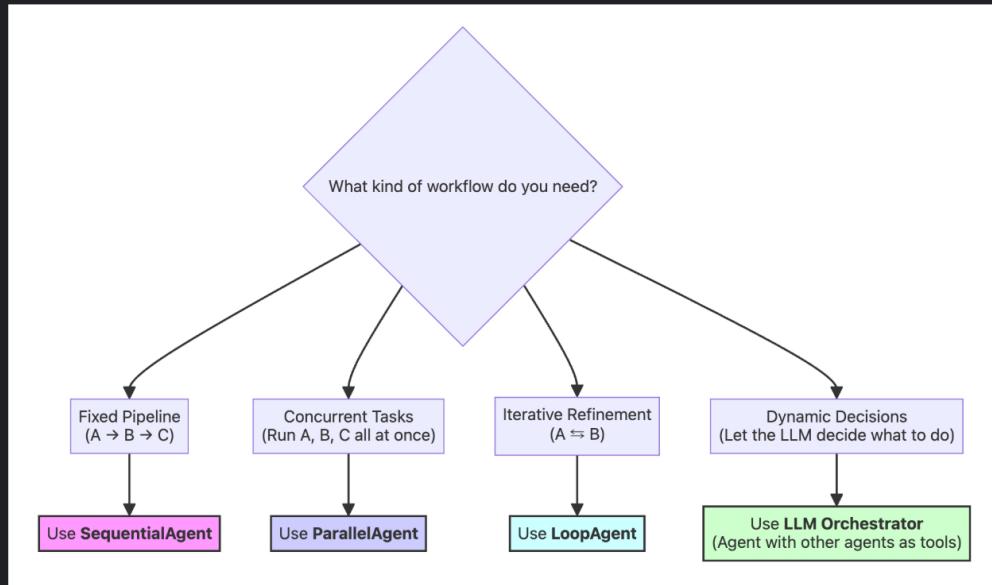
WARNING:google_genai.types:Warning: there are non-text parts in the response: ['function_call'], returning concatenated text result from text parts. check the full candidates.content.parts accessor to get the full model response.

```

You've now implemented a loop agent, creating a sophisticated system that can iteratively review and improve its own output. This is a key pattern for ensuring high-quality results.

You now have a complete toolkit of workflow patterns. Let's put it all together and review how to choose the right one for your use case.

## Section 6: Summary - Choosing the Right Pattern



### Quick Reference Table

Pattern	When to Use	Example	Key Feature
<b>LLM-based (sub_agents)</b>	Dynamic orchestration needed	Research + Summarize	LLM decides what to call
<b>Sequential</b>	Order matters, linear pipeline	Outline → Write → Edit	Deterministic order
<b>Parallel</b>	Independent tasks, speed matters	Multi-topic research	Concurrent execution
<b>Loop</b>	Iterative improvement needed	Writer + Critic refinement	Repeated cycles

### ✓ Congratulations! You're Now an Agent Orchestrator

In this notebook, you made the leap from a single agent to a **multi-agent system**.

You saw **why** a team of specialists is easier to build and debug than one "do-it-all" agent. Most importantly, you learned how to be the **director** of that team.

You used `SequentialAgent`, `ParallelAgent`, and `LoopAgent` to create deterministic workflows, and you even used an LLM as a 'manager' to make dynamic decisions. You also mastered the "plumbing" by using `output_key` to pass state between agents and make them collaborative.

#### 💡 Note: No submission required!

This notebook is for your hands-on practice and learning only. You **do not** need to submit it anywhere to complete the course.

### 📚 Learn More

Refer to the following documentation to learn more:

- [Agents in ADK](#)
- [Sequential Agents in ADK](#)
- [Parallel Agents in ADK](#)
- [Loop Agents in ADK](#)
- [Custom Agents in ADK](#)

### 🎯 Next Steps

Ready for the next challenge? Stay tuned for Day 2 notebooks where we'll learn how to create **Custom Functions**, **use MCP Tools** and manage **Long-Running operations**!

#### Authors

Kristopher Overholt